

AUTONOMOUS PRODUCTION CHOKE CONTROLLER

Problem Statement ID - Honeywell Hackathon, Problem 3A

Problem Statement Title - Autonomous Production Choke Controller for a Single Naturally Flowing Oil Well

Theme - Industrial Process Control / Oil & Gas Automation

PS Category - Software

Student Name - Divyansh Gupta (VIT Vellore)

PROCESS UNDERSTANDING & MODEL

Step test: 120 hr open-loop, choke 30->40->55->45->65 %, Ts = 1 hr

Response is first-order (time constant ~5 hr on oil rate), mild noise

Opening choke RAISES oil rate but LOWERS WHP / FLP / BHP

=> the binding safety constraints are LOWER pressure bounds


Model: ARX(1) per output by least squares

$y[k+1] = a.y[k] + (1-a)(b_0 + b_1.u[k])$; free-run fit 1.5-2.4 % of span

Identified steady-state gains (per % choke):

Q +1.82 bbl/hr WHP -1.59 FLP -0.99 BHP -8.38 psi

Model identification: ARX(1) free-run vs. step-test data



CONTROL STRATEGY

Simplified MPC by brute-force candidate search (per brief)

Each hour, enumerate choke +/- up to 5 % (ramp-feasible, fine grid)


Predict each candidate's pressures over a settling horizon + steady state

REJECT any move predicted to breach WHP/FLP/BHP + safety margin

Among safe moves, minimise $(Q_{ss} - \text{target})^2 + \lambda \cdot (\text{delta choke})^2$

Infeasible target => same cost settles at the highest safe rate

Monotonic first-order response: steady-state feasibility proves the whole trajectory safe -> safe-by-construction, 0 violations



RESULTS - TARGET TRACKING (A & B)

Scenario A - Startup -> 130 bbl/hr

settles 129.9 bbl/hr, steady-state error -0.11, 0 violations



Scenario B - Target step 100 -> 150 bbl/hr

settles 149.7 bbl/hr, steady-state error -0.34, 0 violations

Pressures stay inside the envelope throughout

Choke moves are smooth and respect the +/-5 %/step ramp limit

Feasible-target tracking: Scenario A (startup) and Scenario B (100->150)



SAFETY PERFORMANCE & LESSONS LEARNED

Scenario C - infeasible 185 bbl/hr: controller REFUSES,
settles at max safe 154.9 bbl/hr with 0 violations

Lessons learned:

Binding constraint is a lower pressure bound, not an upper rate cap
Monotonicity makes a cheap brute-force MPC provably safe
Safety margin trades ~4 bbl/hr for zero measured violations


0

MPC violations (A/B/C)

168

naive-baseline breaches on C

Safety-first MPC vs. naive baseline -- Scenario C -- Infeasible Target (185 bbl/hr, settle at max safe)



ARTIFACTS & REFERENCES

Code (flat Python modules at repo root):

dataio, model_id, envelope, simulator, controller, plotting, scenarios, main

Executed Jupyter notebook + technical report.md + results/ (plots, metrics.json)

Run: `pip install -r requirements.txt && python main.py`

main.py asserts 0 envelope violations across A/B/C


Simulator stand-in shares the official .step(choke)->(Q,WHP,FLP,BHP)

signature -> the real event simulator drops in with a one-line change

Repo: github.com/Divyansh2602/autonomous-choke-controller

References: least-squares system identification (Ljung); simplified

MPC by brute-force candidate evaluation (Honeywell brief)



Required 6-trend plot (Scenario C shown)

Technical Report — Autonomous Production Choke Controller (Problem 3A)

Honeywell Hackathon · Single naturally flowing oil well · Control interval $T_s = 1$ hr

1. Problem and approach

Set the production choke (0–100 %) once per hour to track a target oil rate while never violating the pressure operating envelope (lower bounds on wellhead, flowline and bottomhole pressure) or the ± 5 %/step choke ramp limit. If a target is physically unsafe, line out at the maximum achievable safe rate rather than chase it.

Our solution is a **model-based predictive controller that is safe by construction**. We identify a first-order dynamic model of the well from the provided step test, then, every interval, search the reachable choke moves, predict each one's effect on all pressures, and **only ever command a move we have proven keeps the well inside the envelope**. The same decision rule tracks reachable targets tightly and refuses unreachable ones gracefully.

2. Data and open-loop analysis

The dataset is a 120-hour open-loop step test with the choke stepped $30 \rightarrow 40 \rightarrow 55 \rightarrow 45 \rightarrow 65$ %. Two structural facts drive the whole design:

- The response is **first-order** with a time constant of roughly 5 hr on oil rate (slower on pressures), with mild measurement noise.
- **Opening the choke raises oil rate but lowers all three pressures**. The safety constraints are therefore **lower bounds**, and the well is pushed toward them precisely when we chase higher production — the core tension the controller must manage.

3. Dynamic model identification

For each output $y \in \{Q, WHP, FLP, BHP\}$ we fit an ARX(1) model relating the choke u to the output one step ahead:

$$y[k+1] = a y[k] + (1-a) (b_0 + b_1 u[k])$$

This is linear in the parameters, so we recover (a, b_0, b_1) by ordinary least squares (regress $y[k+1]$ on $[y[k], 1, u[k]]$) and read off the physically meaningful quantities: steady-state gain $b_1 = dy_{ss}/du$, intercept b_0 , and time constant $\tau = -Ts/\ln a$.

Output	gain b_1	b_0	τ (hr)	noise σ	free-run RMSE
Q (bbl/hr)	+1.815	39.31	5.15	0.74	1.5 % of span
WHP (psi)	-1.586	319.39	8.52	0.67	1.8 % of span
FLP (psi)	-0.986	219.27	6.76	0.52	2.4 % of span
BHP (psi)	-8.383	3417.35	12.96	2.95	1.9 % of span

The signs confirm the physics (choke up $\rightarrow$ rate up, pressures down). Validation is a genuine **free-run**: the model is initialised once and simulated open-loop under the recorded choke sequence using only its own predictions — never the measured output — and still matches the data to 1.5–2.4 % of signal span ([results/model_validation.png](#)). The one-step residual standard deviations set the noise level of the plant simulator so closed-loop tests see realistic scatter.

4. Operating envelope

The brief specifies no numeric envelope, so we adopt a **documented assumption**, chosen so the maximum safe rate sits high in the observed range and Scenario C is unambiguously infeasible. All limits live in one module ([envelope.py](#)):

Constraint	Value
WHP_min	215 psi
FLP_min	150 psi
BHP_min	2850 psi
Choke range	0–100 %
Ramp limit	± 5 %/step

With the identified maps, **WHP is the first binding constraint**: it reaches 215 psi at choke ≈ 65.8 %,

giving a hard maximum safe rate of $\approx 159 \text{ bbl/hr}$. When the official envelope/simulator is released, only this one file changes.

Per the brief, **wellhead temperature (WHT)** and **annulus pressure (AP)** are part of a complete production operating envelope but are *informational* here — not active constraints in this challenge — so they are not modelled or bounded. The envelope module is structured so they could be added as additional bounded outputs without touching the controller.

5. Controller design

A simplified MPC by brute-force candidate evaluation (explicitly permitted by the brief). Each interval:

1. **Read** current Q, WHP, FLP, BHP and choke position.
2. **Enumerate** candidate next positions: choke $\pm$ up to 5 %, on a 0.5 % grid, clipped to [0, 100] — every candidate already respects the ramp limit.
3. **Predict & screen for safety**. For each candidate, simulate the pressure trajectory over a settling horizon *and* compute its steady state with the ARX models. **Reject** the candidate if any pressure would fall below its bound plus a safety margin, at any predicted step or at steady state.
4. **Optimise**. Among the safe candidates, minimise $(Q_{ss} - \text{target})^2 + \lambda \cdot (\Delta \text{choke})^2$.
5. **Infeasible target / recovery**. Because every safe candidate then has $Q_{ss} < \text{target}$, the cost automatically selects the **highest safe rate**. If no candidate is safe (e.g. the plant starts outside the envelope), the controller backs the choke off to the move that best restores the pressures.

Why this is safe by construction. A stable first-order response moves *monotonically* from the current output to its steady state. So if the current pressure is inside the envelope and the candidate's steady-state pressure is inside it (with margin), every intermediate value is too. Proving steady-state feasibility therefore proves trajectory feasibility — the controller never commits to a move it has not shown to be safe. The safety margin is sized to the fitted process noise ($\approx 4\sigma$ on the binding pressure), which is what turns "predicted-safe" into "measured 0 violations" despite sensor scatter.

6. Results

Three required scenarios, run closed-loop against the simulator (`python main.py`, reproduced in `results/metrics.json`):

Scenario	Target	Settled rate	Settled choke	Steady-state error	Min WHP	Violations
A Startup → Target	130	129.9	50.0 %	−0.11 bbl/hr	238.8 psi	0
B Target Tracking	100 → 150	149.7	61.0 %	−0.34 bbl/hr	221.9 psi	0
C Infeasible	185	154.9	64.0 %	(infeasible)	216.7 psi	0

- **A** brings the well from startup (≈ 90 bbl/hr) to 130 and holds it, tracking to 0.11 bbl/hr with 18.8 psi of pressure headroom to spare.
- **B** tracks a mid-run target step $100 \rightarrow 150$, settling to 0.34 bbl/hr while WHP lines out at 221.9 psi — inside the envelope throughout.
- **C** requests 185 bbl/hr, which would require choke ≈ 80 % and drive WHP to ≈ 192 psi (well below the 215 floor). The controller **refuses**, settling at the maximum safe rate (154.9 bbl/hr) with WHP held at 216.7 psi and **zero violations**.

Each scenario produces the required six-trend plot (Target Q, Actual Q, WHP, FLP, BHP, Choke) in `results/scenario_{A,B,C}.png`.

7. Safety performance vs. a naive baseline

To quantify the value of the safety logic, the same scenarios were run with a naive proportional controller that chases the oil-rate target with no envelope awareness:

Scenario	Safety-first MPC violations	Naive baseline violations
A	0	0
B	0	0
C	0	168 (WHP 57, FLP 56, BHP 55)

On the feasible targets (A, B) both reach the setpoint. On the infeasible target (C) the naive controller drives the choke past 100 %, holding the well ~ 65 psi below the WHP floor for the entire run — 168 constraint breaches — to gain just ~ 30 bbl/hr of unsafe, unsustainable production. The safety-first controller gives that up by design. `results/baseline_comparison_C.png` shows the contrast.

8. Lessons learned & limitations

- **The binding constraint is a lower pressure bound, not an upper rate bound** — recognising that inverted the intuition and made WHP the design driver.
- **Monotonicity is the key that makes a cheap MPC provably safe.** A short brute-force search plus a steady-state feasibility check is enough to guarantee the envelope; no heavy optimiser is required.
- **The safety margin trades a little production for a lot of robustness.** At $\sim 4\sigma$, measured violations are zero at the cost of settling ~ 4 bbl/hr below the theoretical hard limit — a deliberate, adjustable choice.
- **Limitations.** The plant model is first-order and linear. The brief's simulator assumptions (single naturally flowing well, no gas lift/ESP, constant reservoir properties, constant GOR and water cut) keep that valid for this challenge, but a real well would add mild nonlinearity, gas-coning and slugging the ARX model will not capture; the envelope numbers are also our assumption. The controller design is agnostic to all of this: `simulator.WellSimulator` shares the official `.step()` signature, so the real simulator drops in with a one-line change, and the envelope is a single editable module.

9. Reproducing

```
pip install -r requirements.txt
python main.py          # prints model params, runs A/B/C, asserts 0 violations
```

All figures and `metrics.json` are regenerated in `results/`. The narrative walk-through is in `notebook/choke_control_solution.ipynb`.

Figure 1 - ARX(1) model identification: free-run vs. step-test data




Figure 2 - Scenario A (Startup to 130 bbl/hr): six-trend plot




Figure 3 - Scenario B (Target tracking 100 to 150 bbl/hr): six-trend plot




Figure 4 - Scenario C (Infeasible 185 bbl/hr, settle at max safe): six-trend plot




Figure 5 - Safety-first MPC vs. naive baseline on Scenario C

Safety-first MPC vs. naive baseline -- Scenario C -- Infeasible Target (185 bbl/hr, settle at max safe)




Figure 6 - Steady-state maps and the safe operating region

